

# VGG16\_Quantization\_Aware\_Training

December 14, 2025

```
[1]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

from models import *

global best_prec
use_gpu = torch.cuda.is_available()
device = torch.device("cuda" if use_gpu else "mps")    #for mac users
print('=> Building model...')

batch_size = 128
model_name = "vgg16_customized"
model = VGG16_quant()

print(model)

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243,
↪0.262])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
```

```

train=True,
download=True,
transform=transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    normalize,
]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
↳shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
↳shuffle=False, num_workers=2)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch
↳includes "batch_size" data points
# CIFAR10 has 50,000 training data, and 10,000 validation data.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter()
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time()
    for i, (input, target) in enumerate(trainloader):
        # measure data loading time
        data_time.update(time.time() - end)

        input, target = input.to(device), target.to(device)

        # compute output
        output = model(input)
        loss = criterion(output, target)

```

```

        # measure accuracy and record loss
        prec = accuracy(output, target)[0]
        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

    if i % print_freq == 0:
        print('Epoch: [{0}] [{1}/{2}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                epoch, i, len(trainloader), batch_time=batch_time,
                data_time=data_time, loss=losses, top1=top1))

def validate(val_loader, model, criterion ):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

            input, target = input.to(device), target.to(device)

            # compute output
            output = model(input)
            loss = criterion(output, target)

            # measure accuracy and record loss
            prec = accuracy(output, target)[0]

```

```

        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

        if i % print_freq == 0: # This line shows how frequently print out
            the status. e.g., i%5 => every 5 batch, prints out
                print('Test: [{0}/{1}]\t'
                      'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                      'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                      'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                        i, len(val_loader), batch_time=batch_time, loss=losses,
                        top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg

def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True)
    pred = pred.t()
    correct = pred.eq(target.view(1, -1).expand_as(pred))

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0)
        res.append(correct_k.mul_(100.0 / batch_size))
    return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

```

```

def update(self, val, n=1):
    self.val = val
    self.sum += val * n
    self.count += n
    self.avg = self.sum / self.count

def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

def adjust_learning_rate(optimizer, epoch):
    """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120_
    ↪epochs"""
    adjust_list = [60, 100, 120]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.2
    if epoch == 100:
        param_group['momentum'] = param_group['momentum']*0.5

#model = nn.DataParallel(model).cuda()
#all_params = checkpoint['state_dict']
#model.load_state_dict(all_params, strict=False)
#criterion = nn.CrossEntropyLoss().cuda()
#validate(testloader, model, criterion)

```

=> Building model...

```

VGG_quant(
    (features): Sequential(
      (0): QuantConv2d(
        3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): QuantConv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (5): ReLU(inplace=True)

```

```

(6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(7): QuantConv2d(
  64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(9): ReLU(inplace=True)
(10): QuantConv2d(
  128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(12): ReLU(inplace=True)
(13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(14): QuantConv2d(
  128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(16): ReLU(inplace=True)
(17): QuantConv2d(
  256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): QuantConv2d(
  256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(22): ReLU(inplace=True)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(24): QuantConv2d(
  256, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(25): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(26): ReLU(inplace=True)

```

```

(27): QuantConv2d(
  8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(28): ReLU(inplace=True)
(29): QuantConv2d(
  8, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(30): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(31): ReLU(inplace=True)
(32): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(33): QuantConv2d(
  512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(34): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(35): ReLU(inplace=True)
(36): QuantConv2d(
  512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(38): ReLU(inplace=True)
(39): QuantConv2d(
  512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
(40): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(41): ReLU(inplace=True)
(42): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(43): AvgPool2d(kernel_size=1, stride=1, padding=0)
)
(classifier): Linear(in_features=512, out_features=10, bias=True)
)
Files already downloaded and verified
Files already downloaded and verified

```

[2]: # HW

```
# 1. Train with 4 bits for both weight and activation to achieve >90% accuracy
```

```
# 2. Find  $x_{int}$  and  $w_{int}$  for the 2nd convolution layer
# 3. Check the recovered psum has similar value to the un-quantized original
    ↪ psum
# (such as example 1 in W3S2)
```

```
[4]: PATH = "result/vgg16_customized/model_best.pth.tar"
checkpoint = torch.load(PATH)
model.load_state_dict(checkpoint['state_dict'])
use_gpu = torch.cuda.is_available()
device = torch.device("cuda" if use_gpu else "mps")

model.to(device)
model.eval()

test_loss = 0
correct = 0

with torch.no_grad():
    for data, target in testloader:
        data, target = data.to(device), target.to(device) # loading to GPU
        output = model(data)
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(testloader.dataset)

print('\nTest set: Accuracy: {}/{} ({:.3f}%) \n'.format(
    correct, len(testloader.dataset),
    100. * correct / len(testloader.dataset)))
```

Test set: Accuracy: 9182/10000 (91.820%)

```
[5]: #send an input and grap the value by using prehook like HW3
```

```
[6]: class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in)
    def clear(self):
        self.outputs = []

##### Save inputs from selected layer #####
save_output = SaveOutput()
i = 0
```



```

for layer in model.modules():
    i = i+1
    if isinstance(layer, QuantConv2d):
        print(i, "-th layer prehooked")
        layer.register_forward_pre_hook(save_output)
#####

dataiter = iter(testloader)
images, labels = next(dataiter)
images = images.to(device)
out = model(images)

```

```

3 -th layer prehooked
7 -th layer prehooked
12 -th layer prehooked
16 -th layer prehooked
21 -th layer prehooked
25 -th layer prehooked
29 -th layer prehooked
34 -th layer prehooked
38 -th layer prehooked
41 -th layer prehooked
46 -th layer prehooked
50 -th layer prehooked
54 -th layer prehooked

```

```

[7]: w_bit = 4
weight_q = model.features[27].weight_q # quantized value is stored during the
    ↪ training
w_alpha = model.features[27].weight_quant.wgt_alpha # alpha is defined in
    ↪ your model already. bring it out here
w_delta = w_alpha/(2**(w_bit-1)-1) # delta can be calculated by using alpha
    ↪ and w_bit
weight_int = weight_q/w_delta # w_int can be calculated by weight_q and w_delta
print(weight_int) # you should see clean integer numbers

```

```

tensor([[[[ 1.0000,  1.0000, -1.0000],
           [ 2.0000, -0.0000, -1.0000],
           [ 1.0000, -1.0000, -2.0000]],

        [[ 6.0000, -3.0000,  2.0000],
           [ 4.0000, -1.0000, -1.0000],
           [ 3.0000, -7.0000,  5.0000]],

        [[ 3.0000,  2.0000,  2.0000],
           [ 2.0000,  4.0000,  7.0000],

```

```

[ 7.0000,  7.0000,  7.0000]],

[[ 7.0000,  7.0000,  7.0000],
 [ 7.0000,  7.0000,  7.0000],
 [ 7.0000, -3.0000,  1.0000]],

[[ 0.0000, -2.0000,  4.0000],
 [ 0.0000,  0.0000,  2.0000],
 [ 3.0000, -2.0000,  7.0000]],

[[-7.0000,  0.0000,  7.0000],
 [-7.0000,  2.0000,  7.0000],
 [-7.0000,  3.0000,  7.0000]],

[[-1.0000,  1.0000,  1.0000],
 [-2.0000, -0.0000,  1.0000],
 [-2.0000, -1.0000,  2.0000]],

[[ 3.0000, -3.0000, -5.0000],
 [ 2.0000, -2.0000, -5.0000],
 [-3.0000, -5.0000, -6.0000]]],

[[[ 1.0000,  0.0000, -1.0000],
 [ 1.0000,  0.0000, -1.0000],
 [ 1.0000,  0.0000, -1.0000]],

[[-7.0000, -0.0000,  1.0000],
 [-7.0000,  4.0000, -7.0000],
 [-4.0000, -2.0000, -3.0000]],

[[-0.0000, -6.0000, -1.0000],
 [-1.0000,  3.0000,  7.0000],
 [-0.0000,  4.0000,  7.0000]],

[[ 3.0000,  5.0000,  6.0000],
 [ 1.0000,  5.0000,  4.0000],
 [-0.0000, -2.0000,  0.0000]],

[[-4.0000, -1.0000, -0.0000],
 [-3.0000, -1.0000, -5.0000],
 [-3.0000, -4.0000, -7.0000]],

[[ 1.0000,  5.0000,  7.0000],
 [ 5.0000,  3.0000,  5.0000],
 [ 0.0000,  2.0000,  2.0000]],

[[-0.0000, -0.0000,  0.0000],

```

```

[ 1.0000,  1.0000,  1.0000],
[-7.0000, -3.0000,  1.0000]],

[[ 5.0000,  5.0000, -3.0000],
 [ 2.0000,  4.0000, -2.0000],
 [ 0.0000, -1.0000, -2.0000]]],

[[[ 1.0000, -0.0000, -0.0000],
 [ 1.0000, -1.0000,  0.0000],
 [ 0.0000, -1.0000,  0.0000]],

 [[-1.0000,  0.0000, -7.0000],
 [-2.0000,  3.0000, -4.0000],
 [-1.0000, -1.0000, -2.0000]],

 [[ 1.0000,  3.0000,  7.0000],
 [ 0.0000,  7.0000, -3.0000],
 [ 6.0000,  7.0000,  4.0000]],

 [[ 7.0000,  5.0000, -3.0000],
 [ 7.0000,  6.0000, -1.0000],
 [ 2.0000,  7.0000, -4.0000]],

 [[-4.0000, -1.0000, -3.0000],
 [-1.0000, -2.0000, -2.0000],
 [-0.0000, -0.0000, -1.0000]],

 [[-4.0000,  7.0000,  7.0000],
 [-7.0000, -0.0000,  2.0000],
 [-6.0000,  3.0000, -4.0000]],

 [[-1.0000, -0.0000,  1.0000],
 [-2.0000, -0.0000,  1.0000],
 [-7.0000, -2.0000,  1.0000]],

 [[-2.0000, -1.0000, -7.0000],
 [-2.0000,  0.0000, -5.0000],
 [-4.0000, -2.0000,  0.0000]]],

[[[ 1.0000, -0.0000,  0.0000],
 [ 1.0000, -2.0000, -1.0000],
 [ 2.0000, -0.0000, -2.0000]],

 [[-4.0000, -5.0000, -3.0000],
 [-3.0000, -6.0000, -6.0000],
 [ 0.0000, -2.0000, -3.0000]],

```

```

[[-1.0000, -2.0000,  0.0000],
 [-2.0000, -3.0000, -3.0000],
 [-1.0000, -3.0000, -2.0000]],

[[ 0.0000,  0.0000,  1.0000],
 [-1.0000, -2.0000, -2.0000],
 [-1.0000, -2.0000, -4.0000]],

[[ 2.0000, -1.0000, -1.0000],
 [-0.0000, -1.0000, -1.0000],
 [-1.0000,  0.0000,  1.0000]],

[[-1.0000, -1.0000, -1.0000],
 [-1.0000, -3.0000, -2.0000],
 [ 1.0000, -2.0000, -3.0000]],

[[ 1.0000,  0.0000,  0.0000],
 [-2.0000,  0.0000,  0.0000],
 [-1.0000,  0.0000,  1.0000]],

[[ 0.0000, -1.0000, -0.0000],
 [ 1.0000, -3.0000, -3.0000],
 [ 0.0000, -2.0000, -2.0000]]],

[[[ 1.0000,  0.0000,  0.0000],
   [ 2.0000, -0.0000,  1.0000],
   [ 1.0000,  0.0000,  1.0000]],

[[ 2.0000,  6.0000,  7.0000],
 [ 7.0000,  7.0000,  7.0000],
 [-2.0000,  2.0000,  6.0000]],

[[ 0.0000,  3.0000,  1.0000],
 [ 0.0000,  0.0000, -1.0000],
 [ 2.0000,  1.0000,  4.0000]],

[[-1.0000,  1.0000,  4.0000],
 [-2.0000,  0.0000,  1.0000],
 [ 1.0000,  4.0000,  3.0000]],

[[-6.0000, -4.0000, -6.0000],
 [-7.0000, -7.0000, -7.0000],
 [-7.0000, -7.0000, -7.0000]],

[[-1.0000,  0.0000,  4.0000],
 [-2.0000, -0.0000, -3.0000],

```

```

[ 1.0000,  0.0000,  5.0000]],

[[ 0.0000,  0.0000, -0.0000],
 [ 1.0000,  1.0000,  1.0000],
 [-1.0000,  1.0000,  1.0000]],

[[-7.0000, -7.0000, -3.0000],
 [-6.0000,  2.0000,  7.0000],
 [-5.0000, -4.0000, -1.0000]]],

[[[ 1.0000, -0.0000, -0.0000],
 [ 1.0000,  1.0000, -0.0000],
 [ 1.0000,  1.0000, -0.0000]],

 [[-1.0000, -4.0000, -5.0000],
 [-1.0000, -3.0000, -3.0000],
 [-0.0000, -1.0000, -1.0000]],

 [[-0.0000, -1.0000, -1.0000],
 [ 0.0000, -0.0000, -0.0000],
 [-1.0000,  0.0000,  1.0000]],

 [[-0.0000, -1.0000, -0.0000],
 [-0.0000, -2.0000, -1.0000],
 [-0.0000, -0.0000,  0.0000]],

 [[ 1.0000, -2.0000, -0.0000],
 [ 1.0000, -2.0000, -1.0000],
 [ 2.0000, -2.0000, -1.0000]],

 [[ 0.0000, -1.0000, -1.0000],
 [ 0.0000, -3.0000, -3.0000],
 [ 1.0000, -1.0000, -0.0000]],

 [[ 1.0000, -0.0000,  1.0000],
 [ 1.0000,  0.0000,  1.0000],
 [ 4.0000, -0.0000,  1.0000]],

 [[ 0.0000, -2.0000, -1.0000],
 [ 0.0000, -2.0000, -3.0000],
 [ 1.0000, -2.0000, -2.0000]]],

[[[ 1.0000, -0.0000, -0.0000],
 [ 1.0000, -0.0000,  0.0000],
 [ 2.0000, -1.0000,  0.0000]],

```

```

[[-4.0000, -1.0000, -2.0000],
 [-3.0000, -1.0000,  2.0000],
 [ 0.0000, -2.0000, -0.0000]],

[[-2.0000, -1.0000,  7.0000],
 [-1.0000,  3.0000,  5.0000],
 [-1.0000, -0.0000, -0.0000]],

[[ 1.0000,  2.0000,  2.0000],
 [-2.0000,  0.0000,  3.0000],
 [ 1.0000, -1.0000,  2.0000]],

[[ 2.0000,  2.0000,  2.0000],
 [ 2.0000,  0.0000,  2.0000],
 [-3.0000, -5.0000, -3.0000]],

[[ 0.0000,  4.0000,  7.0000],
 [-1.0000, -2.0000,  0.0000],
 [ 0.0000, -3.0000, -4.0000]],

[[ 1.0000,  1.0000,  1.0000],
 [ 0.0000,  1.0000,  2.0000],
 [-3.0000,  2.0000,  1.0000]],

[[ 2.0000, -1.0000, -0.0000],
 [ 1.0000, -2.0000, -1.0000],
 [ 3.0000,  1.0000, -0.0000]],

[[[ 1.0000,  1.0000,  1.0000],
 [ 1.0000,  1.0000,  0.0000],
 [ 1.0000,  1.0000,  1.0000]],

[[-5.0000, -4.0000, -7.0000],
 [ 1.0000,  2.0000, -7.0000],
 [-5.0000, -2.0000, -2.0000]],

[[ 1.0000, -1.0000, -1.0000],
 [ 3.0000, -2.0000, -2.0000],
 [ 2.0000,  1.0000,  0.0000]],

[[ 4.0000,  2.0000, -2.0000],
 [ 4.0000,  2.0000,  2.0000],
 [-1.0000,  1.0000, -2.0000]],

[[ 7.0000,  2.0000, -2.0000],
 [ 7.0000,  7.0000,  6.0000],
 [-7.0000, -0.0000, -2.0000]],

```

```

[[[-2.0000, -1.0000, -4.0000],
  [-4.0000, -0.0000, -1.0000],
  [-1.0000, -1.0000, -4.0000]],

 [[-0.0000, 0.0000, 1.0000],
  [ 1.0000, 1.0000, 2.0000],
  [-2.0000, -1.0000, 2.0000]],

 [[ 2.0000, 0.0000, -1.0000],
  [ 6.0000, 3.0000, -1.0000],
  [ 0.0000, -0.0000, -2.0000]]], device='mps:0',
grad_fn=<DivBackward0>)

```

```

[8]: x_bit = 4
x = save_output.outputs[8][0] # input of the 2nd conv layer
x_alpha = model.features[27].act_alpha
x_delta = x_alpha/(2**x_bit-1)

act_quant_fn = act_quantization(x_bit) # define the quantization function
x_q = act_quant_fn(x, x_alpha) # create the quantized value for x

x_int = x_q/x_delta
print(x_int) # you should see clean integer numbers

```

```

tensor([[[[ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000]],

 [[ 3.0000, 2.0000, 0.0000, 0.0000],
  [ 5.0000, 5.0000, 1.0000, 0.0000],
  [ 4.0000, 4.0000, 1.0000, 0.0000],
  [ 3.0000, 3.0000, 1.0000, 0.0000]],

 [[ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 0.0000, 0.0000, 0.0000]],

 ...,

 [[ 0.0000, 0.0000, 0.0000, 0.0000],
  [ 0.0000, 5.0000, 7.0000, 0.0000],
  [ 9.0000, 15.0000, 15.0000, 2.0000],
  [ 3.0000, 3.0000, 2.0000, 0.0000]],

 [[ 0.0000, 0.0000, 0.0000, 0.0000],

```

```

[ 0.0000, 0.0000, 0.0000, 0.0000],
[ 0.0000, 0.0000, 0.0000, 0.0000],
[ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 3.0000, 6.0000, 5.0000],
 [ 0.0000, 0.0000, 7.0000, 8.0000],
 [ 0.0000, 0.0000, 1.0000, 4.0000],
 [ 0.0000, 2.0000, 4.0000, 4.0000]]],

[[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 1.0000],
 [ 0.0000, 0.0000, 0.0000, 1.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 5.0000, 6.0000, 4.0000, 3.0000],
 [ 5.0000, 5.0000, 3.0000, 2.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 0.0000, 8.0000, 6.0000],
 [ 0.0000, 0.0000, 4.0000, 4.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

...,

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 1.0000, 6.0000, 11.0000, 8.0000],
 [ 0.0000, 7.0000, 11.0000, 8.0000],
 [ 2.0000, 8.0000, 9.0000, 7.0000],
 [ 4.0000, 10.0000, 10.0000, 6.0000]]],

[[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

```



```
[[ 5.0000, 6.0000, 5.0000, 2.0000],
 [ 4.0000, 5.0000, 4.0000, 2.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 1.0000, 1.0000]],
```

```
[[ 0.0000, 0.0000, 5.0000, 3.0000],
 [ 0.0000, 0.0000, 6.0000, 3.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],
```

...,

```
[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],
```

```
[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],
```

```
[[ 1.0000, 5.0000, 8.0000, 6.0000],
 [ 0.0000, 3.0000, 5.0000, 5.0000],
 [ 1.0000, 6.0000, 7.0000, 3.0000],
 [ 3.0000, 5.0000, 6.0000, 4.0000]]],
```

...,

```
[[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],
```

```
[[ 3.0000, 4.0000, 3.0000, 2.0000],
 [ 7.0000, 7.0000, 5.0000, 3.0000],
 [ 6.0000, 5.0000, 2.0000, 1.0000],
 [ 6.0000, 6.0000, 3.0000, 1.0000]],
```

```
[[ 0.0000, 0.0000, 3.0000, 2.0000],
 [ 0.0000, 0.0000, 3.0000, 3.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 1.0000, 0.0000]],
```

...,

```

[[ 0.0000, 0.0000, 1.0000, 0.0000],
 [ 0.0000, 0.0000, 3.0000, 0.0000],
 [ 0.0000, 1.0000, 6.0000, 0.0000],
 [ 0.0000, 0.0000, 2.0000, 0.0000]],

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 0.0000, 2.0000, 1.0000],
 [ 0.0000, 1.0000, 3.0000, 2.0000],
 [ 0.0000, 2.0000, 3.0000, 2.0000],
 [ 3.0000, 6.0000, 6.0000, 3.0000]]],

[[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 3.0000, 4.0000, 3.0000, 2.0000],
 [ 2.0000, 1.0000, 1.0000, 1.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 1.0000, 0.0000, 0.0000, 1.0000]],

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 8.0000, 4.0000],
 [ 0.0000, 0.0000, 1.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

...,

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000],
 [ 0.0000, 0.0000, 0.0000, 0.0000]],

[[ 0.0000, 3.0000, 6.0000, 4.0000],
 [ 0.0000, 4.0000, 6.0000, 4.0000],
 [ 3.0000, 6.0000, 6.0000, 3.0000],
 [ 1.0000, 5.0000, 5.0000, 3.0000]]],

```

```

[[[ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000]],

 [[ 3.0000,  3.0000,  2.0000,  1.0000],
   [ 3.0000,  3.0000,  1.0000,  1.0000],
   [ 4.0000,  4.0000,  1.0000,  0.0000],
   [ 4.0000,  4.0000,  3.0000,  1.0000]],

 [[ 3.0000,  3.0000,  1.0000,  0.0000],
   [ 4.0000,  6.0000,  3.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000]],

 ...,

 [[ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 2.0000,  5.0000,  5.0000,  1.0000],
   [ 2.0000,  8.0000,  7.0000,  1.0000],
   [ 0.0000,  1.0000,  0.0000,  0.0000]],

 [[ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000],
   [ 0.0000,  0.0000,  0.0000,  0.0000]],

 [[ 0.0000,  1.0000,  4.0000,  4.0000],
   [ 0.0000,  0.0000,  1.0000,  3.0000],
   [ 0.0000,  0.0000,  2.0000,  2.0000],
   [ 0.0000,  1.0000,  3.0000,  2.0000]]], device='mps:0',
grad_fn=<DivBackward0>)

```

```

[9]: conv_int = torch.nn.Conv2d(in_channels = 8, out_channels = 8, kernel_size = 3,
    ↪padding=1, bias = False)
conv_int.weight = torch.nn.parameter.Parameter(weight_int)
output_int = conv_int(x_int)    # output_int can be calculated with conv_int
    ↪and x_int

output_recovered = output_int * x_delta * w_delta # recover with x_delta and
    ↪w_delta
relu = nn.ReLU(inplace=True)
output_recovered = relu(output_recovered)

output_ref = save_output.outputs[9][0]

```

```
[19]: difference = abs(output_ref - output_recovered)
print(difference.mean())  ## It should be small, e.g., 2.3 in my trained model
```

```
tensor(4.7602e-07, device='mps:0', grad_fn=<MeanBackward0>)
```

```
[11]: x_pad = torch.zeros(128, 8, 6, 6).to(device)
x_pad[:, :, 1:5, 1:5] = x_int.to(device)
X = x_pad[0]
X = torch.reshape(X, (X.size(0), -1))

### show this cell partially. The following cells should be printed by students
↳###

tile_id = 0
nij = 200 # just a random number
#X = a_tile[tile_id, :, nij:nij+64] # [tile_num, array row num, time_steps]

bit_precision = 4
file = open('activation.txt', 'w') #write to file
file.write('#time0row7[msb-lsb],time0row6[msb-lst],...,time0row0[msb-lst]#\n')
file.write('#time1row7[msb-lsb],time1row6[msb-lst],...,time1row0[msb-lst]#\n')
file.write('#.....#\n')

for i in range(X.size(1)): # time step
    for j in range(X.size(0)): # row #
        X_bin = '{0:04b}'.format(round(X[7-j,i].item()))
        for k in range(bit_precision):
            file.write(X_bin[k])
            #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
file.close() #close file
```

```
[12]: X.size()
```

```
[12]: torch.Size([8, 36])
```

```
[13]: ### Complete this cell ###
tile_id = 0
kij = 0
W = torch.reshape(weight_int, (weight_int.size(0), weight_int.size(1), -1))
#W = w_tile[tile_id, :, :, kij] # w_tile[tile_num, array col num, array row num,
↳kij]

bit_precision = 4
for kij in range(9):
    name = 'weight_kij'+str(kij)+'.txt'
    file = open(name, 'w') #write to file
```

```

file.write('#col0row7[msb-lsb],col0row6[msb-lst],...,col0row0[msb-lst]#\n')
file.write('#col1row7[msb-lsb],col1row6[msb-lst],...,col1row0[msb-lst]#\n')
file.write('#.....#\n')

for i in range(W.size(0)):
    for j in range(W.size(1)):
        if (W[i, 7-j, kij].item()<0):
            W_bin = '{0:04b}'.format(round(W[i,7-j,kij].
↪item()+2**bit_precision))
        else:
            W_bin = '{0:04b}'.format(round(W[i,7-j,kij].item()))
        for k in range(bit_precision):
            file.write(W_bin[k])
            #file.write(' ') # for visibility with blank between words, you
↪can use
        file.write('\n')
    file.close() #close file

```

[14]: W.size()

[14]: torch.Size([8, 8, 9])

```

[15]: p_nijg = range(X.size(1))
psum = torch.zeros(8, len(p_nijg), 9).to(device)

for kij in range(9):
    for nij in p_nijg: # time domain, sequentially given input
        m = nn.Linear(8, 8, bias=False)
        m.weight = torch.nn.Parameter(W[:, :, kij])
        relu = nn.ReLU(inplace=True)
        psum[:, nij, kij] = relu(m(X[:, nij])).to(device)

bit_precision = 16
for kij in range(9):
    name = 'psum_kij'+str(kij)+'.txt'
    file = open(name, 'w') #write to file
    file.write('#time0col7[msb-lsb],time0col6[msb-lst],....
↪,time0col0[msb-lst]#\n')
    file.write('#time1col7[msb-lsb],time1col6[msb-lst],....
↪,time1col0[msb-lst]#\n')
    file.write('#.....#\n')
    for i in range(psum.size(1)):
        for j in range(psum.size(0)):
            if (psum[7-j,i,kij].item()<0):
                P_bin = '{0:016b}'.format(round(psum[7-j,i,kij].
↪item()+2**bit_precision))
            else:

```

```

        P_bin = '{0:016b}'.format(round(psum[7-j,i,kij].item()))
    for k in range(bit_precision):
        file.write(P_bin[k])
        #file.write(' ') # for visibility with blank between words, you
↪ can use
        file.write('\n')
    file.close()

```

```
[16]: print(psum.size())
```

```
torch.Size([8, 36, 9])
```

```
[17]: out = output_int[0]
out = torch.reshape(out, (out.size(0), -1))

bit_precision = 16
file = open('output.txt', 'w') #write to file
file.write('#time0col7[msb-lsb],time0col6[msb-lst],...,time0col0[msb-lst]#\n')
file.write('#time1col7[msb-lsb],time1col6[msb-lst],...,time1col0[msb-lst]#\n')
file.write('#.....#\n')

for i in range(out.size(1)):
    for j in range(out.size(0)):
        if (out[7-j,i].item()<0):
            O_bin = '{0:016b}'.format(round(out[7-j,i].item()+2**bit_precision))
        else:
            O_bin = '{0:016b}'.format(round(out[7-j,i].item()))
        for k in range(bit_precision):
            file.write(O_bin[k])
            #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
    file.close()

```

```
[18]: print(out.size())
```

```
torch.Size([8, 16])
```

```
[ ]:
```